Date

Prompt engineering

Subject

LLM Generates text, how??

- we give llm a prompt ((a sequence of tokens))
- llm generates text one token at a time.
- each new token is generated conditioned on:
  - the original prompt
  - all tokens generated so far

next-token prediction Formula

- at generation step i, the model computes:
  $P(w_i | w_{j<i})$
- the model outputs a distribution over a vocabulary tokens
- we choose the next token by:
  - ☐ sampling, or ☐ Greed choice: or ☐ beam search
- this repeats untill:
  - ☐ "end of text" token appears
  - ☐ a preset length is reached

Date Prompts Subject

- **prompt**: a text string that user provides to instruct the llm.

- **purpose**:
  - tells the model what task to perform
  - provides context, style, target format, constraints

- **Prompt engineering**:
  - the process of designing prompts that makes the model perform a task well.
  - Good prompts:
    - clearly describe the task
    - specify only constraints (format, style, length)

- **Few-shot vs zero-shot prompting**

  - **zero shot prompting**:
    - Prompts containes instructions but no examples

    example:
    "translate into French"

  - **Few-shot prompting**:
    - we include labeled examples in the prompt
    - Helps the model understand the pattern or task.

    examples: english → persian
    dog → سگ
    house → خونه

Date Context learning Subject

Demonstrations in Prompting:

- Demonstrations: small labelled examples inserted into the prompt, usually drawn from a labeled training set.

- How demonstrations are choosen.
  - sometimes manually selected by humans
  - sometime chosen automatically
    - an optimizer searches for the set of demo examples that gives the best performance.

- number of demonstrations:
  - only a few examples usually needed ("few-shot")
  - adding more gives diminishing returns.
  - too many can cause:
    - overfitting to specific examples
    - worse generalization

- what demonstrations really achieve??
  - main purpose: show the task structure.
  - They do not need to show perfectly. even demonstrations with wrong answers can still help the model.

JAHAN

Date

In - context learning

Subject

- Definition:
  In - context learning is when a language model improves it's behaviour on a task just by being shown examples in the prompt, not by updating its weights.
  - the learning happens on the fly, inside model's temporary context window.

- the learning that happens during prompting, improves the model's performance without:
  - gradient descent
  - weight updates
  - retraining